

SINDHUJA CHADUVULA

Associate Applied ML Specialist

Toronto, Canada | +1 437-799-7401 | schaduvu@uwaterloo.ca | <https://www.linkedin.com/in/sindhuja-chaduvula/>

TECHNICAL SKILLS

Languages and Scripting: Python, C/C++, R, SQL, Javascript, Bash/Shell.

ML/AI Frameworks: Langchain, Langsmith, Langgraph, Phidata, CrewAI, LlamaIndex, OpenAI, OLLama, Hugging Face, Pytorch, Tensorflow, PySpark, Pandas, Numpy, OpenCV, NLTK, MLflow, Weights & Biases.

Generative AI & NLP: RAG, Finetuning LLMs, Prompt Engineering, Computer Vision, Deep Learning, Time Series Forecasting.

Cloud & DevOps: AES(S3), Azure, GCP, Docker, Kubernetes, Github Actions, Linux.

Soft Skills: Communication Skills, Time Management, Leadership, Adaptability, Teamwork, Decision Making, Problem Solving

EDUCATION

University of Waterloo

Jan 2023 – Apr 2024

Master of Engineering in Electrical and Computer Engineering

GPA – **3.90/4.00**

Key Focus: Data Modelling & Analysis, Image Processing, Text Analytics (NLP), Data Analysis & Management

EXPERIENCE

Associate Applied ML Specialist, Vector Institute (Toronto, Canada)

Sept 2025 – Present

- **Generated 125 synthetic text scenes and 125 MCQs** in the hiring and legal domain using GPT-4o Mini to train a bias detection model.
- **Researched AI explainability methods** (intrinsic, post-hoc, concept-based, mechanistic, and prototype-based) as part of an AI Safety project to improve model interpretability.
- Conducted **DPO training** with hint sampling and repeated sampling for **1,000 samples** of the Skywork reward preference dataset using the Qwen2-7B model.

Machine Learning Associate, Vector Institute (Toronto, Canada)

May 2025 – Sept 2025

- Developed a LangChain-based RAG pipeline to parse and semantically index **1,000-page Schedule of Benefits** documents using LlamaParse and Qdrant, **reducing manual look-up time for physicians by over 80%**.
- Optimized document retrieval by engineering hybrid strategies (OpenAIEmbeddings + BM25) with billing code filters, achieving **Recall@5 of 0.9** and **BERTScore of 0.8**, significantly boosting clinical response accuracy.
- Independently built the end-to-end retrieval and generation stack using Azure OpenAI models, **improving response precision and enabling scalable billing code assistance for clinical users**.
- Presented weekly results to **C-level stakeholders** to align delivery with strategic goals in a resource-constrained environment

Data Scientist, Genellipse (Mississauga, Canada)

Aug 2024 – Dec 2024

- Led a 4-member AI team to automate SEC 10-Q filings by fine-tuning LLMs (Gemma 2B, LLaMA 3, Mistral 7B, OpenAI) on **600 historical filings** across 75 public companies, **reducing manual workload by 60%**.
- Built scalable ETL pipelines using the EDGAR API to extract filings, process structured/unstructured data, and prepared LLM-ready datasets in AWS S3, **accelerating document readiness by 70%**.
- **Containerized** Mistral-based submodels using **Docker** for distinct SEC filing sections, and integrated **automated JSON structure validation** to ensure **reliable downstream parsing and deployment**.
- Delivered a functional PoC for SEC 10-Q generation with **40% template-level fill accuracy**, aligned with ethical compliance protocols and highlighting model adaptation gaps for roadmap planning.

AI Engineer, AIFocal (Markham, Canada)

Sept 2023 – Oct 2023

- Built a PDF-aware QA chatbot using Gemma 2B and LangChain RAG, integrating Milvus for retrieval to **enable accurate summarization and long-form reasoning over user-uploaded content**.
- Designed a **prompt routing system** using keyword + semantic similarity for dynamic query handling.
- Achieved **85% accuracy on synthetic QA benchmarks**, validating system performance for extracting answers from complex PDFs.
- Contributed to prompt engineering, retrieval tuning, and vector DB design as part of a 3-member Agile team that **delivered a working demo within 4 weeks**.

Assistant System Engineer, Tata Consultancy Services (Hyderabad, India)

Oct 2021 – Dec 2022

- Built **churn prediction models (Random Forest, XGBoost)** with **87% accuracy** on 5 years of telecom data; visualized key features using **Matplotlib** and **Seaborn**.
- Engineered **preprocessing pipelines** (capping, flooring, outlier handling) to support churn analysis; collaborated in Agile cycles to deliver actionable insights for retention.

TECHNICAL PROJECTS

Multi-PDF Conversational Assistant

Tech: Langchain, FAISS, Open AI, Streamlit, Information Retrieval

- Built an **AI-powered chatbot** to parse and answer questions from multiple PDFs using **vector indexing** and **semantic search**.
- Integrated **FAISS** with **LangChain** for efficient retrieval and deployed the solution with a **Streamlit UI**.